

Este proyecto esta compuesto de los 3 sprints que se realizaron a lo largo del modulo, puede haber ligeras modificaciones de esto a lo que se llevo a cabo en cada sprints por los comentarios asi como los feedbacks recibidos

- Sprint 1: Desarrollar SVM y Softmax: En este caso vi el codigo que puso el profesor en los comentarios y se me hizo mas simple usar ese codigo hecho en clase debido al ahorro de transformar las variables categoricas a numericas, en este caso solo se necesitan ajustar los parametros, es decir, de eliminan las categorias desde un principio y se configura de una manera mas rapida los modelos.
- sprint 2: Se entrenaron 4 arboles de decisiones pero añadi un 5 donde este no tiene ningun hiper parametro
- Sprint 3: Se desarrollo un ensemble asi como un random forest

Cabe resaltar que en cada sprint vuelvo a cargar la datra por las modificaciones que hago en cada uno de ellos

▼ Librerias a usar

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from google.colab import drive
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import LinearSVC
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn import tree
from sklearn.metrics import f1_score
from sklearn.ensemble import BaggingClassifier
from sklearn.metrics import mean_squared_error
from sklearn.ensemble import AdaBoostRegressor
from sklearn.tree import DecisionTreeRegressor
drive.mount('/content/drive')
```

🔗 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.moun



sprint 1

▼ Carga de datos

```
! pip install kaggle
! mkdir ~/.kaggle
! cp /content/drive/MyDrive/kaggle.json ~/.kaggle/kaggle.json
! chmod 600 ~/.kaggle/kaggle.json

! kaggle datasets download parulpandey/palmer-archipelago-antarctica-penguin-data
! unzip palmer-archipelago-antarctica-penguin-data.zip
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.7/dist-packages (1.5.12)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: certifi in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from kaggle)
mkdir: cannot create directory '/root/.kaggle': File exists
palmer-archipelago-antarctica-penguin-data.zip: Skipping, found more recently modified 1
Archive: palmer-archipelago-antarctica-penguin-data.zip
replace penguins_lter.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: penguins_lter.csv
replace penguins_size.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: penguins_size.csv
```



▼ Analisis de datos

```
penguins = pd.read_csv('penguins_size.csv').dropna()
penguins.head()
```

	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0

```
# Esta instrucción ayuda a comprobar si un DataFrame contiene valores nulos en cualquier atributo
penguins.isnull().values.any()
```

```
False
```

▼ Preparacion de datos


```
X = penguins.drop(["species","island","sex"], axis=1)
```

```
y = penguins.species.astype("category").cat.codes
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=.2, random_state=42)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=.2, random_state=42)
```

```
X_train
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	
230	40.9	13.7	214.0	4650.0	
84	37.3	17.8	191.0	3350.0	
303	50.0	15.9	224.0	5350.0	
22	35.9	19.2	189.0	3800.0	
29	40.5	18.9	180.0	3950.0	
...	
194	50.9	19.1	196.0	3550.0	
77	37.2	19.4	184.0	3900.0	
112	39.7	17.7	193.0	3200.0	
277	45.5	15.0	220.0	5000.0	
108	38.1	17.0	181.0	3175.0	

267 rows × 4 columns

▼ SVM

[illegible]

```

ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning:
ConvergenceWarning,
GridSearchCV(cv=5, estimator=LinearSVC(max_iter=10000),
            param_grid=[{'C': [0.1, 0.5, 1, 2, 5],
                          'loss': ['hinge', 'squared_hinge']},
                        {'C': [0.1, 0.5, 1, 2, 5], 'fit_intercept': [False],

```

```
grid_search.best_estimator_
```

```
LinearSVC(C=5, fit_intercept=False, max_iter=10000)
```

```

cvresults = grid_search.cv_results_
for mean_score, params in zip(cvresults["mean_test_score"], cvresults["params"]):
    print(np.sqrt(-mean_score), params)

```

```

0.39298531311519613 {'C': 0.1, 'loss': 'hinge'}
0.38308031026015243 {'C': 0.1, 'loss': 'squared_hinge'}
0.6566330909059309 {'C': 0.5, 'loss': 'hinge'}
0.6753625731199567 {'C': 0.5, 'loss': 'squared_hinge'}
0.541305717505674 {'C': 1, 'loss': 'hinge'}
0.5752889469360648 {'C': 1, 'loss': 'squared_hinge'}
0.707748864799822 {'C': 2, 'loss': 'hinge'}
0.6809784550075622 {'C': 2, 'loss': 'squared_hinge'}
0.4985653772173597 {'C': 5, 'loss': 'hinge'}
0.610071538117657 {'C': 5, 'loss': 'squared_hinge'}
0.5464452177125045 {'C': 0.1, 'fit_intercept': False, 'loss': 'hinge'}
0.6039122394950113 {'C': 0.1, 'fit_intercept': False, 'loss': 'squared_hinge'}
0.4863618700343812 {'C': 0.5, 'fit_intercept': False, 'loss': 'hinge'}
0.7131587395809371 {'C': 0.5, 'fit_intercept': False, 'loss': 'squared_hinge'}
0.6799001036500849 {'C': 1, 'fit_intercept': False, 'loss': 'hinge'}
0.4251032677596799 {'C': 1, 'fit_intercept': False, 'loss': 'squared_hinge'}
0.47103377396152696 {'C': 2, 'fit_intercept': False, 'loss': 'hinge'}
0.6199569426675574 {'C': 2, 'fit_intercept': False, 'loss': 'squared_hinge'}
0.5450366863962642 {'C': 5, 'fit_intercept': False, 'loss': 'hinge'}
0.3578997669551949 {'C': 5, 'fit_intercept': False, 'loss': 'squared_hinge'}

```

```

svm_classifier_best = Pipeline([
    ("scaler", StandardScaler()),
    ("linear_svc", LinearSVC(C=5, loss="hinge", fit_intercept=False, m
])

```

```
svm_classifier_best.fit(X_train, y_train)
```

```

/usr/local/lib/python3.7/dist-packages/sklearn/svm/_base.py:1208: ConvergenceWarning: Li
ConvergenceWarning,
Pipeline(steps=[('scaler', StandardScaler()),
                 ('linear_svc',
                  LinearSVC(C=5, fit_intercept=False, loss='hinge',
                             max_iter=10000))])

```

```

y_pred = svm_classifier_best.predict(X_test)
accuracy_score(y_test, y_pred)

```

```
0.9850746268656716
```

▼ Softmax

```

softmax_reg = LogisticRegression(multi_class="multinomial", solver="lbfgs", max_iter=10000)
softmax_reg.fit(X_train, y_train)

```

```
LogisticRegression(max_iter=10000, multi_class='multinomial')
```

```
score = softmax_reg.score(X, y)
score
```

```
0.9970059880239521
```

Splint 2

▼ carga de datos

```
! cp /content/drive/MyDrive/kaggle.json ~/.kaggle/kaggle.json
! chmod 600 ~/.kaggle/kaggle.json
```

```
! kaggle datasets download parulpandey/palmer-archipelago-antarctica-penguin-data
! unzip palmer-archipelago-antarctica-penguin-data.zip
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.7/dist-packages (1.5.12)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: urllib3 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: certifi in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from kaggle)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from kaggle)
mkdir: cannot create directory '/root/.kaggle': File exists
palmer-archipelago-antarctica-penguin-data.zip: Skipping, found more recently modified 1
Archive: palmer-archipelago-antarctica-penguin-data.zip
replace penguins_lter.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: penguins_lter.csv
replace penguins_size.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: penguins_size.csv
```

▼ Analisis de datos

```
penguins = pd.read_csv('penguins_size.csv').dropna()
penguins.head()
```

	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0

```
# Esta instrucción ayuda a comprobar si un DataFrame contiene valores nulos en cualquier atributo
penguins.isnull().values.any()
```

```
False
```

▼ Preparacion de datos


```
X = penguins.drop(["species","island","sex"], axis=1)
```

```
y = penguins.species.astype("category").cat.codes
```

```
# X_train, X_test, y_train, y_test = train_test_split(
#     X, y, test_size=.2, random_state=42)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=.4, random_state=42)
```

```
X_train
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	
28	37.9	18.6	172.0	3150.0	
42	36.0	18.5	186.0	3100.0	
253	59.6	17.0	230.0	6050.0	
73	45.8	18.9	197.0	4150.0	
273	50.1	15.0	225.0	5000.0	
...	
194	50.9	19.1	196.0	3550.0	
77	37.2	19.4	184.0	3900.0	
112	39.7	17.7	193.0	3200.0	
277	45.5	15.0	220.0	5000.0	
108	38.1	17.0	181.0	3175.0	

200 rows × 4 columns


```
X_nombres = list(X.columns)
especies = penguins.species.unique()
```

▼ Árboles

```
tree1 = DecisionTreeClassifier(max_depth=4)
tree1.fit(X_train, y_train)
tree2 = DecisionTreeClassifier(max_depth=4, min_samples_split=8)
tree2.fit(X_train, y_train)
tree3 = DecisionTreeClassifier(max_depth=4, min_samples_split=8, max_features=3)
tree3.fit(X_train, y_train)
tree4 = DecisionTreeClassifier(min_samples_split=8, max_features=3)
tree4.fit(X_train, y_train)
tree5 = DecisionTreeClassifier()
tree5.fit(X_train, y_train)
```

```
DecisionTreeClassifier()
```

```
print("Predicciones sobre Test")
y_prediction1 = tree1.predict(X_test)
print("Arbol 1 " +str(accuracy_score(y_test, y_prediction1)))
```

```
y_prediction2 = tree2.predict(X_test)
print("Arbol 2 " +str(accuracy_score(y_test, y_prediction2)))
```

```
y_prediction3 = tree3.predict(X_test)
print("Arbol 3 " +str (accuracy_score(y_test, y_prediction3)))
```

```
y_prediction4 = tree4.predict(X_test)
print("Arbol 4 " +str(accuracy_score(y_test, y_prediction4)))
```

```
y_prediction5 = tree5.predict(X_test)
print("Arbol 5 " +str(accuracy_score(y_test, y_prediction5)))
```

```
Predicciones sobre Test
Arbol 1 0.9701492537313433
Arbol 2 0.9552238805970149
Arbol 3 0.9701492537313433
Arbol 4 0.9626865671641791
Arbol 5 0.9626865671641791
```

```
print("Predicciones sobre Training")
y_prediction1t = tree1.predict(X_train)
print("Arbol 1 " +str(accuracy_score(y_train, y_prediction1t)))
```

```
y_prediction2t = tree2.predict(X_train)
print("Arbol 2 " +str(accuracy_score(y_train, y_prediction2t)))
```

```

y_prediction3t = tree3.predict(X_train)
print("Arbol 3 " +str (accuracy_score(y_train, y_prediction3t)))

y_prediction4t = tree4.predict(X_train)
print("Arbol 4 " +str(accuracy_score(y_train, y_prediction4t)))

y_prediction5t = tree5.predict(X_train)
print("Arbol 5 " +str(accuracy_score(y_train, y_prediction5t)))

```

```

Predicciones sobre Trining
Arbol 1 0.985
Arbol 2 0.985
Arbol 3 0.985
Arbol 4 0.995
Arbol 5 1.0

```

```

y_prediction1

```

```

array([0, 2, 0, 1, 0, 2, 2, 1, 1, 1, 0, 0, 2, 0, 2, 0, 0, 1, 0, 2, 0, 0,
       2, 1, 0, 0, 2, 2, 1, 2, 1, 2, 0, 0, 2, 2, 1, 2, 0, 0, 0, 0, 1, 1,
       0, 0, 2, 0, 0, 2, 0, 2, 2, 0, 0, 2, 0, 0, 1, 1, 2, 2, 2, 0, 0, 2,
       0, 2, 0, 2, 0, 0, 0, 2, 2, 1, 0, 0, 0, 2, 0, 0, 0, 0, 2, 2, 0, 2,
       1, 1, 2, 2, 0, 0, 0, 0, 0, 1, 0, 0, 0, 2, 2, 0, 1, 0, 2, 2, 0, 1,
       0, 2, 0, 1, 1, 1, 0, 2, 0, 0, 0, 2, 2, 0, 0, 2, 0, 0, 0, 1, 0, 0,
       1, 0], dtype=int8)

```

```

print(confusion_matrix(y_test,y_prediction1))
print(classification_report(y_test,y_prediction1))

```

```

[[65  0  0]
 [ 2 24  1]
 [ 1  0 41]]

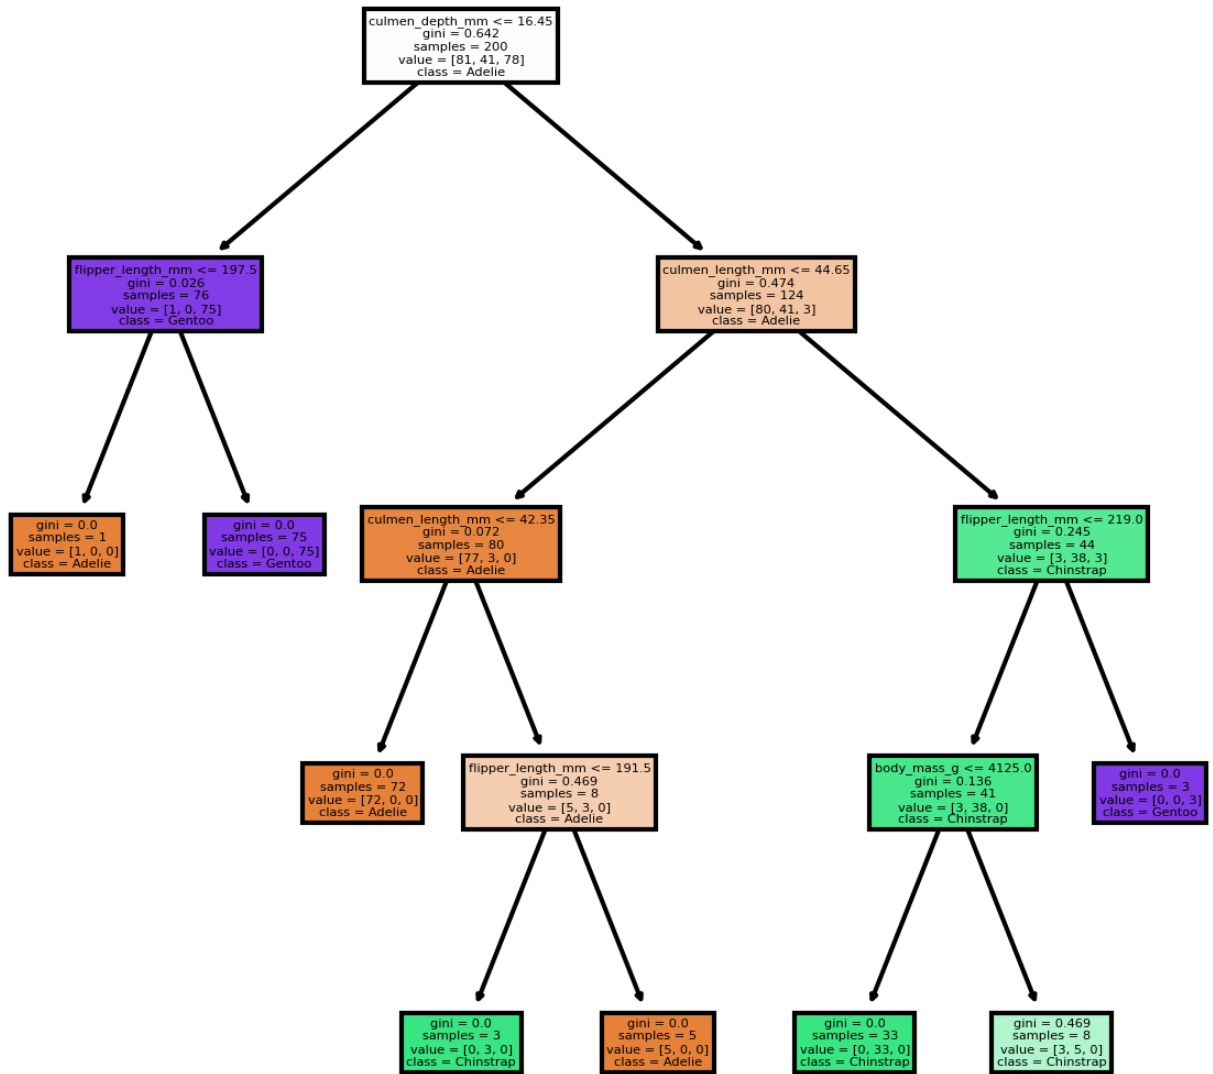
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	65
1	1.00	0.89	0.94	27
2	0.98	0.98	0.98	42
accuracy			0.97	134
macro avg	0.98	0.96	0.96	134
weighted avg	0.97	0.97	0.97	134

```

fig, acex = plt.subplots(nrows =1 ,ncols =1 ,figsize=(5,5), dpi=300)
tree.plot_tree(tree1,
                feature_names = X_nombres,
                class_names = especies,
                filled = True)
fig.show()

```



```
y_prediction2
```

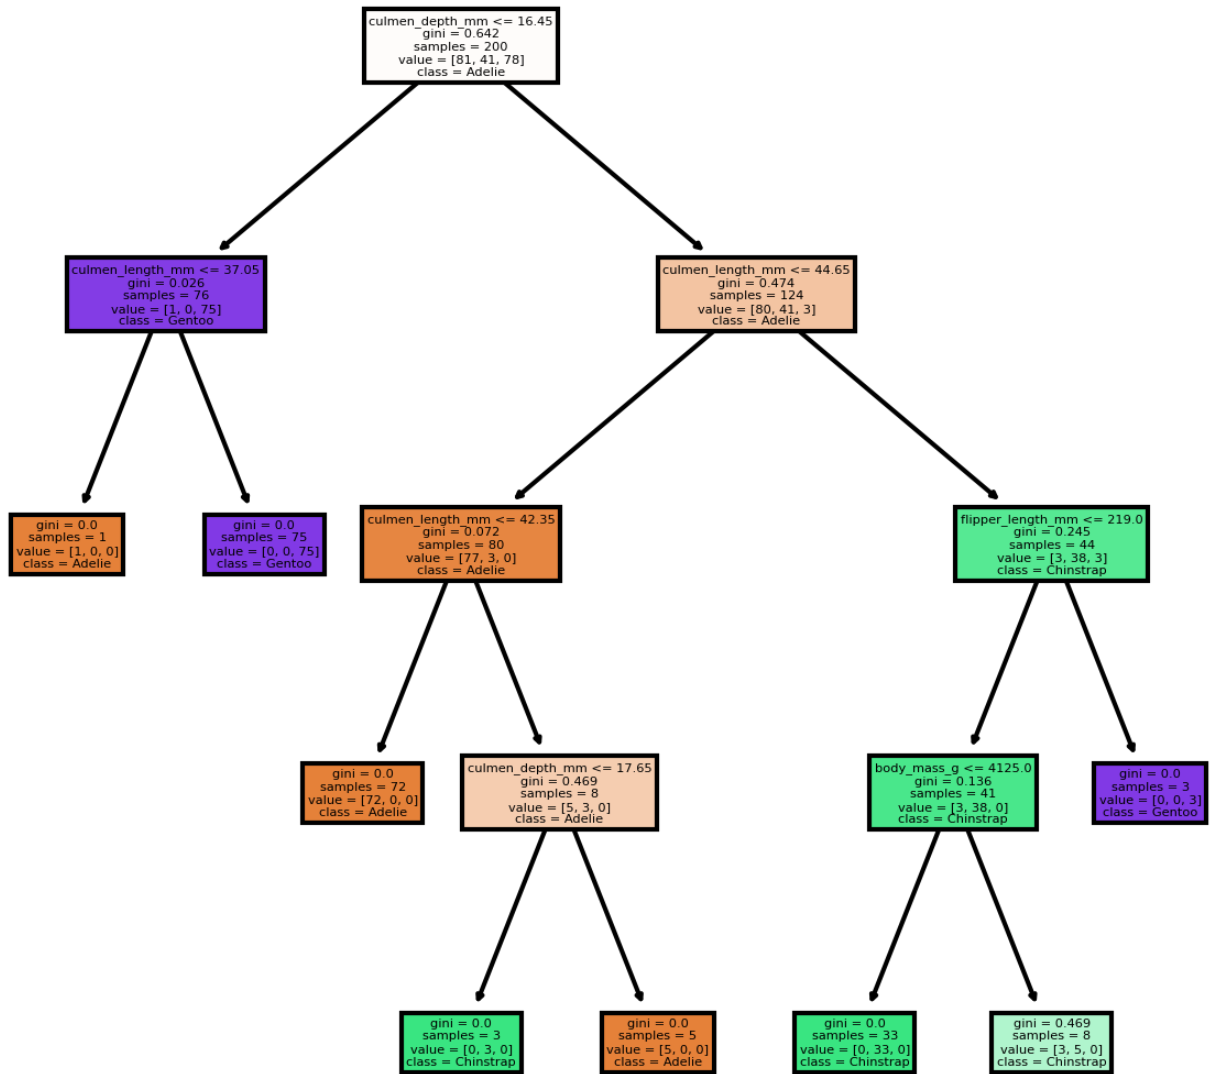
```
array([0, 2, 0, 1, 0, 2, 2, 1, 1, 1, 1, 0, 2, 0, 2, 0, 0, 1, 0, 2, 0, 0,
       2, 1, 0, 0, 2, 2, 1, 2, 1, 2, 0, 0, 2, 2, 1, 2, 0, 0, 0, 0, 1, 1,
       0, 0, 2, 0, 0, 2, 0, 2, 2, 0, 0, 2, 0, 0, 1, 1, 2, 2, 2, 0, 0, 2,
       0, 2, 0, 2, 0, 0, 1, 2, 2, 1, 0, 0, 0, 2, 0, 0, 0, 0, 2, 2, 0, 2,
       1, 1, 2, 2, 0, 0, 0, 0, 0, 1, 0, 0, 0, 2, 2, 0, 1, 0, 2, 2, 0, 1,
       0, 2, 0, 1, 1, 1, 0, 2, 0, 0, 0, 2, 2, 0, 0, 2, 0, 0, 0, 1, 2, 0,
       1, 0], dtype=int8)
```

```
print(confusion_matrix(y_test,y_prediction2))
print(classification_report(y_test,y_prediction2))
```

```
[[63  1  1]
 [ 2 24  1]
 [ 0  1 41]]
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	65
1	0.92	0.89	0.91	27
2	0.95	0.98	0.96	42
accuracy			0.96	134
macro avg	0.95	0.94	0.95	134
weighted avg	0.95	0.96	0.96	134

```
fig, acex = plt.subplots(nrows =1 ,ncols =1 ,figsize=(5,5), dpi=300)
tree.plot_tree(tree2,
                feature_names = X_nombres,
                class_names = especies,
                filled = True)
fig.show()
```



```
y_prediction3
```

```
array([0, 2, 0, 1, 0, 2, 2, 1, 1, 1, 0, 0, 2, 0, 2, 0, 0, 1, 0, 2, 0, 0,
       2, 1, 0, 0, 2, 2, 1, 2, 1, 2, 0, 0, 2, 2, 1, 2, 0, 0, 0, 0, 1, 1,
       0, 0, 2, 0, 0, 2, 0, 2, 2, 0, 0, 2, 0, 0, 1, 1, 2, 2, 2, 0, 0, 2,
       0, 2, 0, 2, 0, 0, 0, 2, 2, 1, 0, 0, 0, 2, 0, 0, 0, 0, 2, 2, 0, 2,
       1, 1, 2, 2, 0, 0, 0, 0, 0, 1, 0, 0, 0, 2, 2, 0, 1, 0, 2, 2, 0, 1,
       0, 2, 0, 1, 1, 1, 0, 2, 0, 0, 0, 0, 2, 0, 0, 2, 0, 0, 0, 1, 0, 0,
       1, 0], dtype=int8)
```

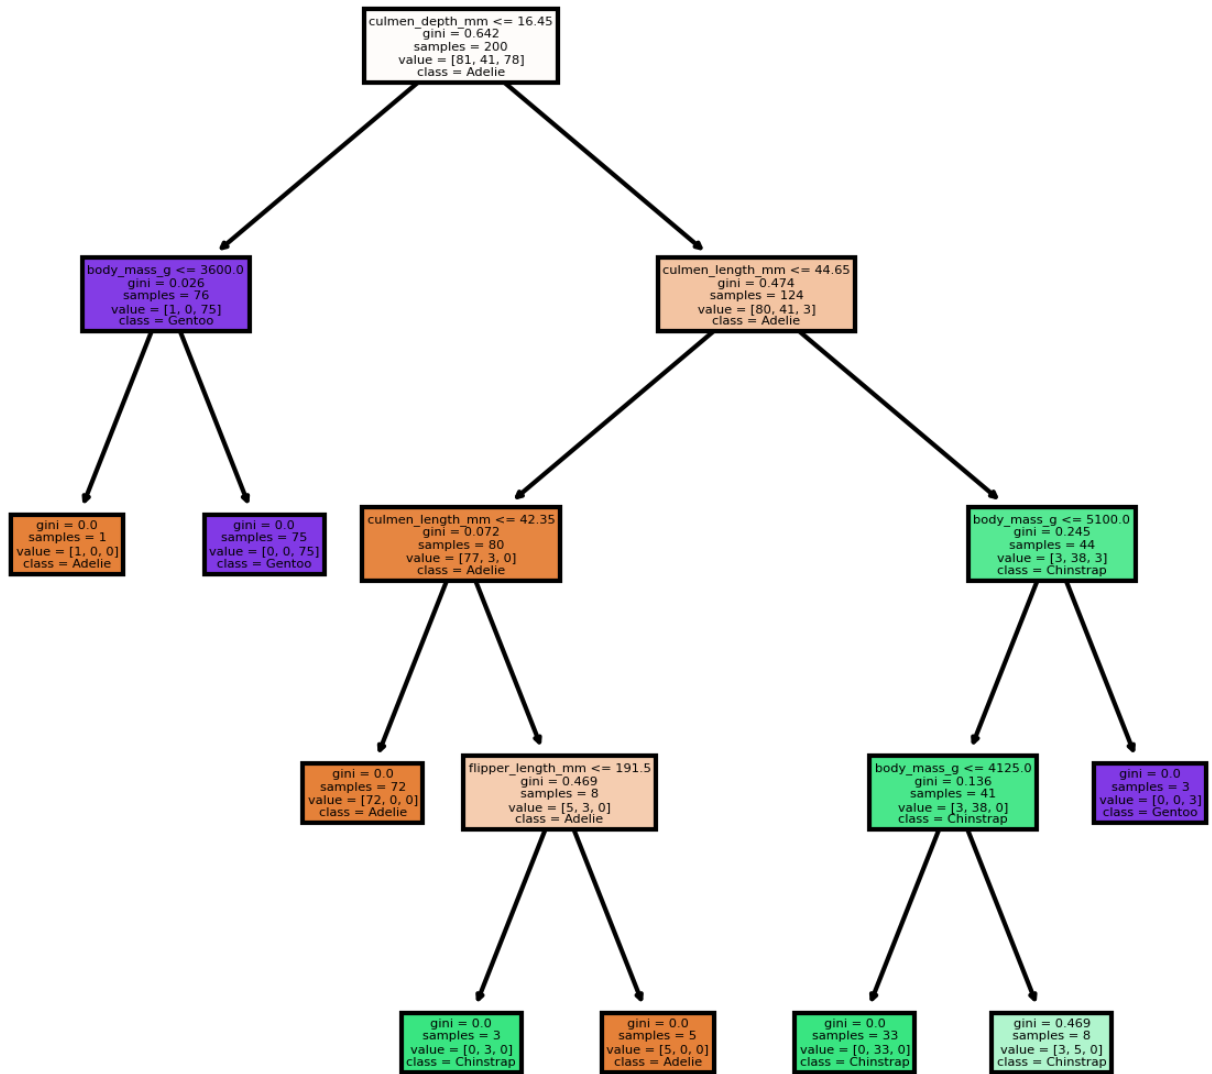
```
# tree.plot_tree(tree3)
```

```
print(confusion_matrix(y_test,y_prediction3))
print(classification_report(y_test,y_prediction3))
```

```
[[65  0  0]
 [ 3 24  0]
 [ 1  0 41]]
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	65
1	1.00	0.89	0.94	27
2	1.00	0.98	0.99	42
accuracy			0.97	134
macro avg	0.98	0.96	0.97	134
weighted avg	0.97	0.97	0.97	134

```
fig, acex = plt.subplots(nrows =1 ,ncols =1 ,figsize=(5,5), dpi=300)
tree.plot_tree(tree3,
                feature_names = X_nombres,
                class_names = especies,
                filled = True)
fig.show()
```



y_prediction4

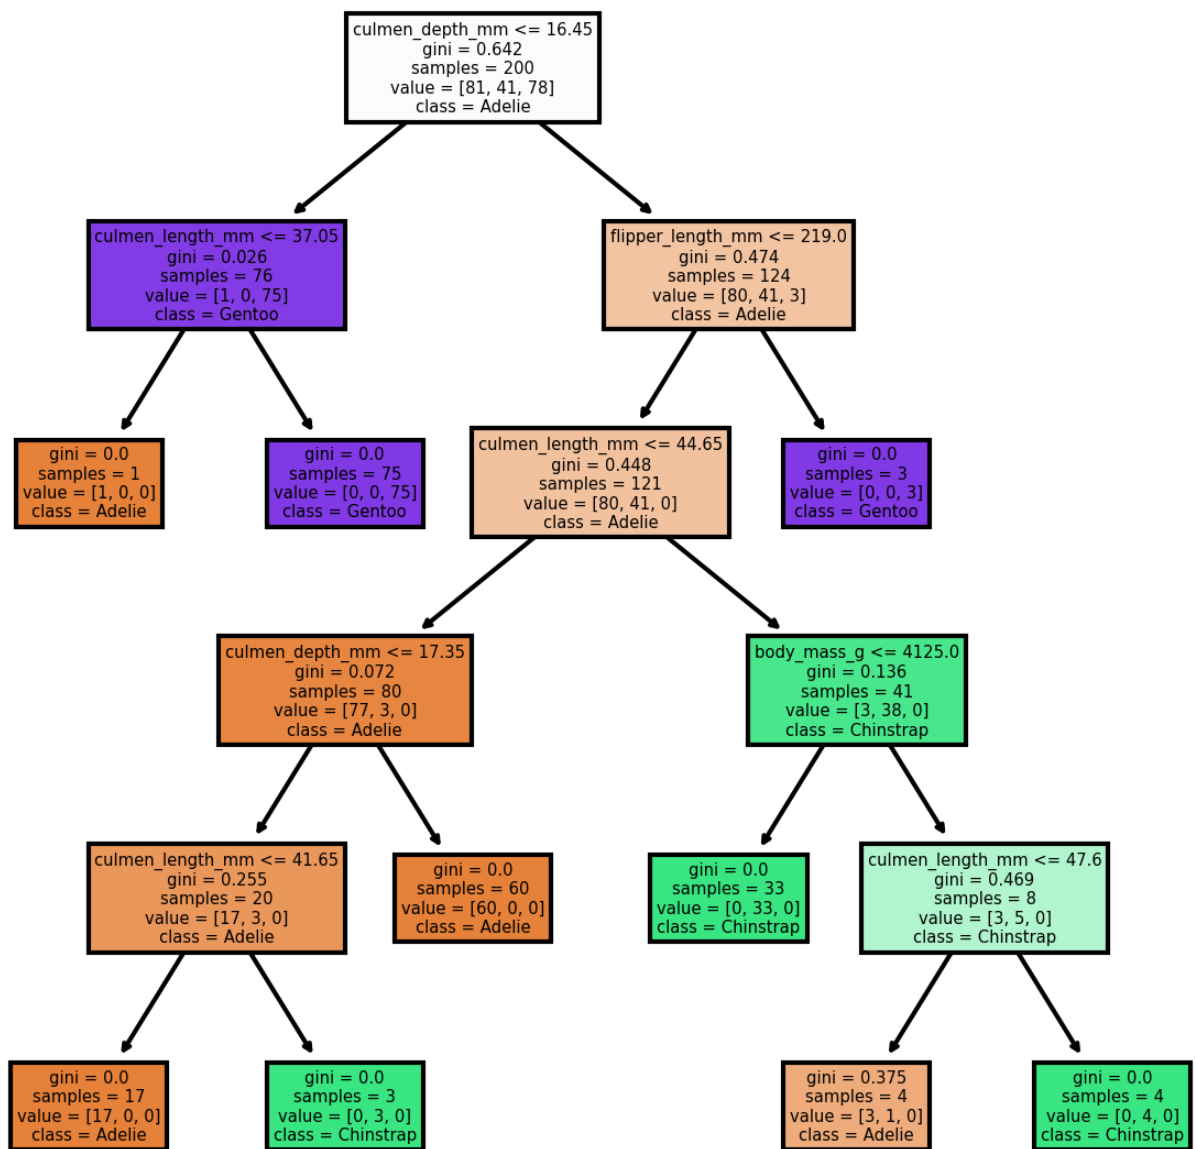
```
array([0, 2, 0, 1, 0, 2, 2, 1, 1, 1, 0, 0, 2, 0, 2, 0, 0, 1, 0, 2, 0, 0,
       2, 1, 0, 0, 2, 2, 1, 2, 1, 2, 0, 0, 2, 2, 1, 2, 0, 0, 0, 0, 1, 1,
       0, 0, 2, 0, 0, 2, 0, 2, 2, 0, 0, 2, 0, 0, 1, 1, 2, 2, 2, 0, 0, 2,
       0, 2, 0, 2, 0, 0, 1, 2, 2, 1, 0, 0, 0, 2, 0, 0, 0, 0, 2, 2, 0, 2,
       1, 1, 2, 2, 0, 0, 0, 0, 0, 1, 0, 0, 0, 2, 2, 0, 1, 0, 2, 2, 0, 1,
       0, 2, 0, 1, 1, 1, 0, 2, 0, 0, 0, 2, 2, 0, 0, 2, 0, 0, 0, 1, 2, 0,
       1, 0], dtype=int8)
```

```
print(confusion_matrix(y_test,y_prediction4))
print(classification_report(y_test,y_prediction4))
```

```
[[64  0  1]
 [ 2 24  1]
 [ 0  1 41]]
```

	precision	recall	f1-score	support
0	0.97	0.98	0.98	65
1	0.96	0.89	0.92	27
2	0.95	0.98	0.96	42
accuracy			0.96	134
macro avg	0.96	0.95	0.95	134
weighted avg	0.96	0.96	0.96	134

```
fig, axes = plt.subplots(nrows =1 ,ncols =1 ,figsize=(5,5), dpi=300)
tree.plot_tree(tree4,
                feature_names = X_nombres,
                class_names = especies,
                filled = True)
fig.show()
```

```
y_prediction5
```

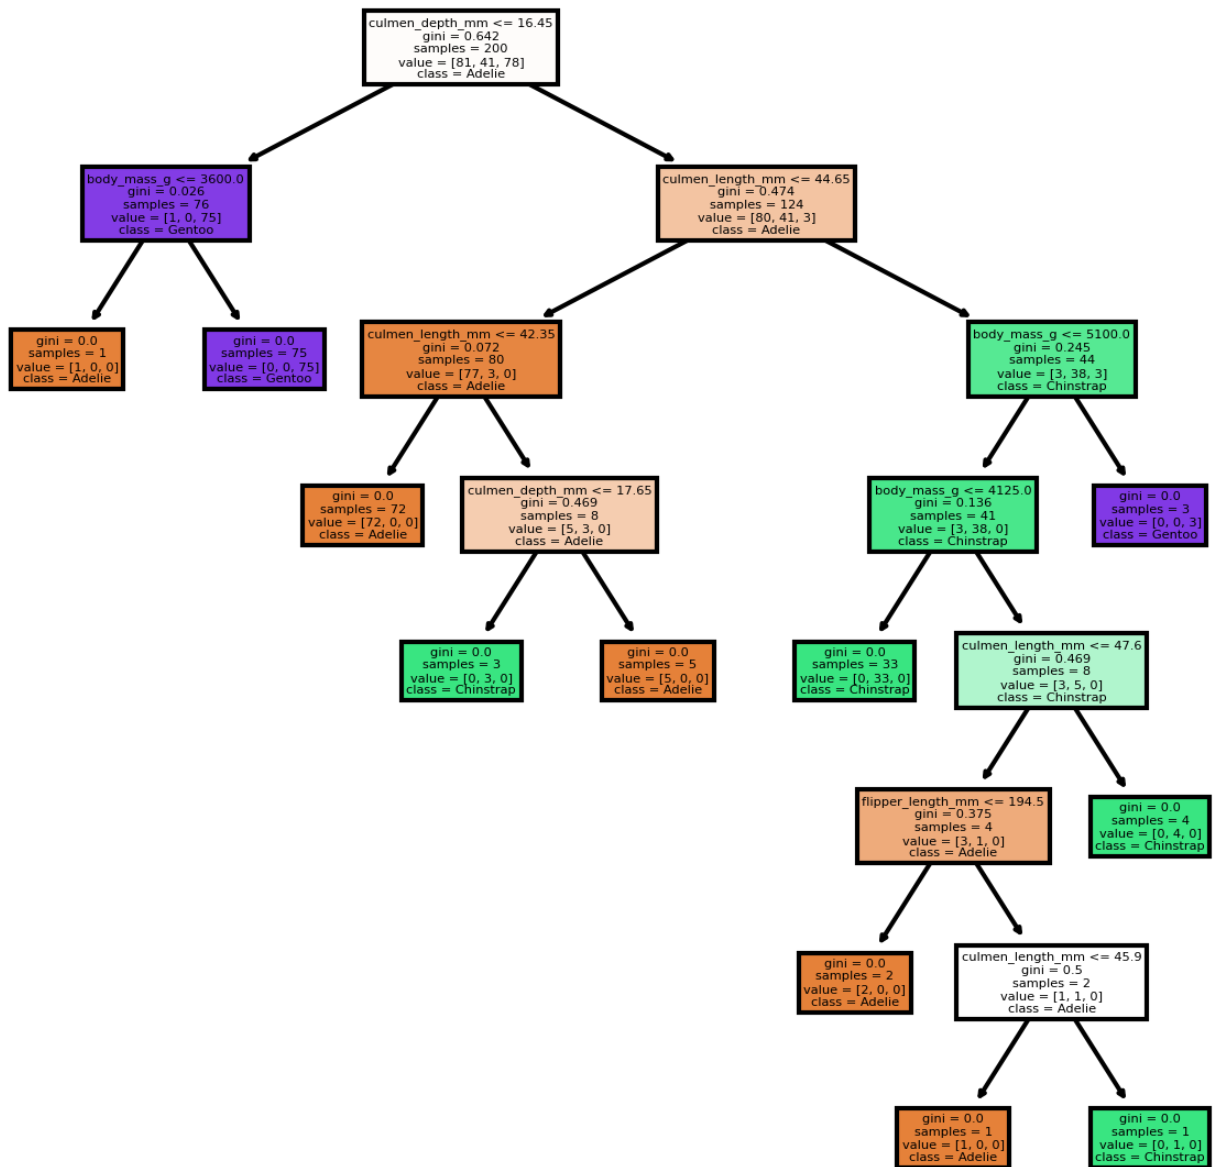
```
array([0, 2, 0, 1, 0, 2, 2, 1, 1, 1, 1, 0, 2, 0, 2, 0, 0, 1, 0, 2, 0, 0,
       2, 1, 0, 0, 2, 2, 1, 2, 1, 2, 0, 0, 2, 2, 1, 2, 0, 0, 0, 0, 1, 1,
       0, 0, 2, 0, 0, 2, 0, 2, 2, 0, 0, 2, 0, 0, 1, 1, 2, 2, 2, 0, 0, 2,
       0, 2, 0, 2, 0, 0, 1, 2, 2, 1, 0, 0, 0, 2, 0, 0, 0, 0, 2, 2, 0, 2,
       1, 1, 2, 2, 0, 0, 0, 0, 0, 1, 0, 0, 0, 2, 2, 0, 1, 0, 2, 2, 0, 1,
       0, 2, 0, 1, 1, 1, 0, 2, 0, 0, 0, 0, 2, 0, 0, 2, 0, 0, 0, 1, 0, 0,
       1, 0], dtype=int8)
```

```
print(confusion_matrix(y_test,y_prediction5))
print(classification_report(y_test,y_prediction5))
```

```
[[64  1  0]
 [ 3 24  0]
 [ 0  1 41]]
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	65
1	0.92	0.89	0.91	27
2	1.00	0.98	0.99	42
accuracy			0.96	134
macro avg	0.96	0.95	0.95	134
weighted avg	0.96	0.96	0.96	134

```
fig, axes = plt.subplots(nrows =1 ,ncols =1 ,figsize=(5,5), dpi=300)
tree.plot_tree(tree5,
                feature_names = X_nombres,
                class_names = especies,
                filled = True)
fig.show()
```



Sprint 3

▼ Carga de datos

```
! cp /content/drive/MyDrive/kaggle.json ~/.kaggle/kaggle.json
! chmod 600 ~/.kaggle/kaggle.json
! kaggle datasets download parulpandey/palmer-archipelago-antarctica-penguin-data
! unzip palmer-archipelago-antarctica-penguin-data.zip


palmer-archipelago-antarctica-penguin-data.zip: Skipping, found more recently modified ]
Archive:  palmer-archipelago-antarctica-penguin-data.zip
replace penguins_lter.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: penguins_lter.csv
replace penguins_size.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: penguins_size.csv
```

▼ Analisis de datos

```
import pandas as pd
penguins = pd.read_csv('penguins_size.csv')
penguins.head()
```

	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0
3	Adelie	Torgersen	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0

```
penguins.describe()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	
count	342.000000	342.000000	342.000000	342.000000	

```
penguins.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   species               344 non-null   object
1   island                344 non-null   object
2   culmen_length_mm      342 non-null   float64
3   culmen_depth_mm       342 non-null   float64
4   flipper_length_mm     342 non-null   float64
5   body_mass_g           342 non-null   float64
6   sex                   334 non-null   object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB
```

▼ Transformacion de datos

```
penguins.drop('sex', axis=1,inplace=True)
df_gender = pd.get_dummies(penguins['island'])
penguins.drop('island', axis=1,inplace=True)
penguins = pd.concat([penguins, df_gender], axis=1)
penguins
```

```

species culmen_length_mm culmen_depth_mm flipper_length_mm body_mass_g Biscoe
penguins.isna().any()

species      False
culmen_length_mm  True
culmen_depth_mm  True
flipper_length_mm  True
body_mass_g      True
Biscoe          False
Dream           False
Torgersen       False
dtype: bool

```

```

penguins['culmen_length_mm'].fillna((penguins['culmen_length_mm'].mean()), inplace=True)
penguins['culmen_depth_mm'].fillna((penguins['culmen_depth_mm'].mean()), inplace=True)
penguins['flipper_length_mm'].fillna((penguins['flipper_length_mm'].mean()), inplace=True)
penguins['body_mass_g'].fillna((penguins['body_mass_g'].mean()), inplace=True)

```

```

343    Gentoo      49.9      16.1      213.0      5400.0      1

```

```

from sklearn.model_selection import train_test_split

```

```

X = penguins.drop("species", axis=1)

```

```

y = penguins.species.astype("category").cat.codes

```

```

# Usar esta instrucción cuando se requieran etiquetas de clase codificadas numéricamente, p.e
# 0 si vamos a hacer búsqueda de hiperparámetros para optimizar modelos.

```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=.2, random_state=42)


```

▼ transformacion de categoricas

```

X_train_num = X_train.drop(["Biscoe", "Dream", "Torgersen"], axis=1) # Obtener una versión solo
from sklearn.impute import SimpleImputer
num_imputer = SimpleImputer(strategy="median") # Rellenar valores perdidos de atributos numér
X_train_num_array = num_imputer.fit_transform(X_train_num)
X_train_num = pd.DataFrame(X_train_num_array, columns=X_train_num.columns, index=X_train_num.
X_train_num.head()

```


	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	
66	35.5	16.2	195.0	3350.0	

```
from sklearn.impute import SimpleImputer
```

```
num_imputer = SimpleImputer(strategy="median") # Rellenar valores perdidos de atributos numér
```

```
X_train_num_array = num_imputer.fit_transform(X_train_num)
```

```
X_train_num = pd.DataFrame(X_train_num_array, columns=X_train_num.columns, index=X_train_num.  
X_train_num.head()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	
66	35.5	16.2	195.0	3350.0	
229	46.8	15.4	215.0	5150.0	
7	39.2	19.6	195.0	4675.0	
140	40.2	17.1	193.0	3400.0	
323	49.1	15.0	228.0	5500.0	


```
X_train_num = X_train.drop(["Biscoe", "Dream", "Torgersen"], axis=1) # Obtener una versión solo
```

```
from sklearn.impute import SimpleImputer
```

```
num_imputer = SimpleImputer(strategy="median") # Rellenar valores perdidos de atributos numér
```

```
X_train_num_array = num_imputer.fit_transform(X_train_num)
```

```
X_train_num = pd.DataFrame(X_train_num_array, columns=X_train_num.columns, index=X_train_num.  
X_train_num.head()
```


	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	
66	35.5	16.2	195.0	3350.0	
229	46.8	15.4	215.0	5150.0	
7	39.2	19.6	195.0	4675.0	
140	40.2	17.1	193.0	3400.0	
323	49.1	15.0	228.0	5500.0	

```
from sklearn.impute import SimpleImputer
```

```
num_imputer = SimpleImputer(strategy="median") # Rellenar valores perdidos de atributos numér
```

```
X_train_num_array = num_imputer.fit_transform(X_train_num)
```

```
X_train_num = pd.DataFrame(X_train_num_array, columns=X_train_num.columns, index=X_train_num.  
X_train_num.head()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	
66	35.5	16.2	195.0	3350.0	
229	46.8	15.4	215.0	5150.0	
7	39.2	19.6	195.0	4675.0	
140	40.2	17.1	193.0	3400.0	
323	49.1	15.0	228.0	5500.0	

▼ division de conjunto

```
X = penguins.drop('species', axis=1)
X
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	Biscoe	Dream
0	39.10000	18.70000	181.000000	3750.000000	0	0
1	39.50000	17.40000	186.000000	3800.000000	0	0
2	40.30000	18.00000	195.000000	3250.000000	0	0
3	43.92193	17.15117	200.915205	4201.754386	0	0
4	36.70000	19.30000	193.000000	3450.000000	0	0
...
339	43.92193	17.15117	200.915205	4201.754386	1	0
340	46.80000	14.30000	215.000000	4850.000000	1	0
341	50.40000	15.70000	222.000000	5750.000000	1	0
342	45.20000	14.80000	212.000000	5200.000000	1	0
343	49.90000	16.10000	213.000000	5400.000000	1	0

344 rows × 7 columns

```
y = penguins.species
y
```

```
0    Adelie
1    Adelie
2    Adelie
3    Adelie
4    Adelie
...
339  Gentoo
340  Gentoo
```



```

341     Gentoo
342     Gentoo
343     Gentoo
Name: species, Length: 344, dtype: object

```

▼ 4. Random Forests

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size = 0.4, random_state=42)

scaler = StandardScaler()

scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
# Las dos anteriores instrucciones serían equivalentes a:
# X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

bagging_clf = BaggingClassifier(
    DecisionTreeClassifier(), n_estimators=200,
    max_samples=100, bootstrap=True, n_jobs=-1)

bagging_clf.fit(X_train, y_train)

y_pred10 = bagging_clf.predict(X_test[:10])
y_pred10

array(['Chinstrap', 'Chinstrap', 'Gentoo', 'Chinstrap', 'Gentoo',
       'Gentoo', 'Gentoo', 'Gentoo', 'Adelie', 'Gentoo'], dtype=object)

y_test[:10]

```

```

194     Chinstrap
157     Chinstrap
225         Gentoo
208     Chinstrap
318         Gentoo
329         Gentoo
319         Gentoo
260         Gentoo
114         Adelie
220         Gentoo
Name: species, dtype: object

```

```
y_pred = bagging_clf.predict(X_test)
```

```
print(classification_report(y_pred, y_test)) # Una visión más detallada del rendimiento por c
```

	precision	recall	f1-score	support
Adelie	0.95	0.98	0.97	62
Chinstrap	0.96	0.92	0.94	25
Gentoo	1.00	0.98	0.99	51
accuracy			0.97	138
macro avg	0.97	0.96	0.97	138
weighted avg	0.97	0.97	0.97	138

```
confusion_matrix(y_pred, y_test)
```

```
array([[61,  1,  0],  
       [ 2, 23,  0],  
       [ 1,  0, 50]])
```

```
accuracy_score(y_pred, y_test)
```

```
0.9710144927536232
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size = 0.2, random_state=42)
```

```
scaler = StandardScaler()
```

```
scaler.fit(X_train)
```

```
X_train_scaled = scaler.transform(X_train)
```

```
# Las dos anteriores instrucciones serían equivalentes a:
```

```
# X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
bagging_clf = BaggingClassifier(  
    DecisionTreeClassifier(), n_estimators=200,  
    max_samples=100, bootstrap=True, n_jobs=-1)
```

```
bagging_clf.fit(X_train, y_train)
```

```
y_pred10 = bagging_clf.predict(X_test[:10])
```

```
y_pred10
```

```
array(['Chinstrap', 'Chinstrap', 'Gentoo', 'Chinstrap', 'Gentoo',
```

```
'Gentoo', 'Gentoo', 'Gentoo', 'Adelie', 'Gentoo'], dtype=object)
```

```
y_test[:10]
```

```
194    Chinstrap
157    Chinstrap
225      Gentoo
208    Chinstrap
318      Gentoo
329      Gentoo
319      Gentoo
260      Gentoo
114      Adelie
220      Gentoo
Name: species, dtype: object
```

```
y_pred = bagging_clf.predict(X_test)
```

```
BaggingClassifier(base_estimator=DecisionTreeClassifier(), max_samples=100,
                  n_estimators=200, n_jobs=-1)
```

```
BaggingClassifier(base_estimator=DecisionTreeClassifier(), max_samples=100,
                  n_estimators=200, n_jobs=-1)
```

```
print(classification_report(y_pred, y_test)) # Una visión más detallada del rendimiento por c
```

	precision	recall	f1-score	support
Adelie	1.00	0.97	0.98	33
Chinstrap	0.94	1.00	0.97	15
Gentoo	1.00	1.00	1.00	21
accuracy			0.99	69
macro avg	0.98	0.99	0.98	69
weighted avg	0.99	0.99	0.99	69

```
confusion_matrix(y_pred, y_test)
```

```
array([[32,  1,  0],
       [ 0, 15,  0],
       [ 0,  0, 21]])
```

```
accuracy_score(y_pred, y_test)
```

```
0.9855072463768116
```

```
df_num = penguins.drop('species', axis=1) # Eliminar variable categórica
```

```

df_num.head()

X = df_num.drop('body_mass_g',axis=1)
y = df_num['body_mass_g']

# Particionado en entrenamiento + test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
ada_reg = AdaBoostRegressor(
    DecisionTreeRegressor(max_depth=2), n_estimators=100,
    learning_rate=0.5)

ada_reg.fit(X_train, y_train)

    AdaBoostRegressor(base_estimator=DecisionTreeRegressor(max_depth=2),
        learning_rate=0.5, n_estimators=100)

y_pred = ada_reg.predict(X_test)

mse = mean_squared_error(y_pred, y_test)
rmse = np.sqrt(mse)
rmse

    407.0114669170131

penguins.body_mass_g.mean()/397.39171939550465

    10.573331503626804

df_num = penguins.drop('species', axis=1) # Eliminar variable categórica
df_num.head()

X = df_num.drop('body_mass_g',axis=1)
y = df_num['body_mass_g']

# Particionado en entrenamiento + test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.4, random_state = 42)
ada_reg = AdaBoostRegressor(
    DecisionTreeRegressor(max_depth=2), n_estimators=100,
    learning_rate=0.5)

ada_reg.fit(X_train, y_train)

    AdaBoostRegressor(base_estimator=DecisionTreeRegressor(max_depth=2),
        learning_rate=0.5, n_estimators=100)

y_pred = ada_reg.predict(X_test)

mse = mean_squared_error(y_pred, y_test)

```

```
rmse = np.sqrt(mse)
rmse
```

```
373.7301469260958
```

```
penguins.body_mass_g.mean()/377.55171876274784
```

```
11.128950491165105
```

✓ 0 s completado a las 16:39

